

BigQuery Schema Design Mastery: From Fundamentals to Enterprise Optimization

Data Engineering Interview Preparation

August 17, 2025

Contents

1	The Evolution of Data Warehouse Schema Design	3
1.1	Traditional Relational Database Limitations	3
1.2	The BigQuery Revolution	3
2	Real-World Case Study: GO-JEK's Data Architecture	3
2.1	Business Context and Scale	3
2.2	Data Structure Complexity	4
3	Schema Design Approaches: Comprehensive Analysis	4
3.1	Approach 1: Normalized (Relational) Schema	4
3.2	Approach 2: Fully Denormalized Schema	5
3.3	Approach 3: Nested and Repeated Fields (BigQuery Optimal)	6
4	BigQuery Data Types Deep Dive	7
4.1	STRUCT Data Type: Nested Fields	7
4.2	ARRAY Data Type: Repeated Fields	8
5	Schema Design Decision Framework	9
5.1	When to Use Nested and Repeated Fields	9
5.2	When to Keep Normalized Schemas	10
6	Performance Optimization Strategies	10
6.1	Partitioning: Foundation of Performance	10
6.2	Clustering: Advanced Performance Tuning	11
6.3	Query Optimization Techniques	13
7	Real-World Implementation Examples	13
7.1	E-commerce Platform Schema	13
7.2	Analytics Platform Schema	15
8	Interview Questions and Discussion Points	16
8.1	Core Concept Questions	16
8.2	Performance Optimization Questions	17
8.3	Real-World Scenario Questions	18

9	Advanced Topics and Best Practices	18
9.1	Schema Design Patterns	18
9.2	Performance Monitoring and Optimization	19
10	Key Takeaways and Best Practices	19
10.1	Schema Design Principles	19
10.2	Performance Optimization	19
10.3	Cost Management	20
10.4	Implementation Guidelines	20

Executive Summary

This comprehensive guide covers BigQuery schema design from fundamental concepts to enterprise-level optimization strategies. We explore the paradigm shift from traditional relational databases to BigQuery's nested and repeated fields, using real-world examples like GO-JEK's 13+ petabyte monthly processing to illustrate practical applications.

1 The Evolution of Data Warehouse Schema Design

1.1 Traditional Relational Database Limitations

Historical Context: Traditional relational databases were designed for transactional processing with normalized schemas that prioritize data integrity and storage efficiency over analytical query performance.

Key Limitations in Analytics Context:

- **JOIN Performance:** Most intensive computational workloads in analytical queries
- **Record-based Processing:** Must open entire records to extract JOIN keys
- **Scalability Issues:** 10+ table JOINS become prohibitively expensive
- **Query Complexity:** Need to know all related tables in advance
- **Storage Model Mismatch:** Row-based storage inefficient for column-based analytics

1.2 The BigQuery Revolution

Paradigm Shift: BigQuery introduced columnar storage with support for nested and repeated fields, fundamentally changing how we approach data warehouse schema design.

Key Innovations:

- **Columnar Storage:** Only reads required columns, ignoring unused data
- **Nested Fields (STRUCT):** Pre-joined data within single tables
- **Repeated Fields (ARRAY):** Multiple values within single rows
- **Serverless Architecture:** Automatic scaling and optimization
- **Pay-per-query Model:** Cost optimization through efficient processing

2 Real-World Case Study: GO-JEK's Data Architecture

2.1 Business Context and Scale

Company Profile:

- **Service:** Indonesia's leading ride-booking and multi-service platform

- **Data Volume:** 13+ petabytes processed monthly
- **Use Case:** Real-time business decision support
- **Challenge:** Efficient storage and querying of complex ride-booking data

2.2 Data Structure Complexity

Core Data Entities:

- **Orders:** Single pickup and drop-off locations per ride
- **Events:** Multiple events per order (ordered, confirmed, en route, completed)
- **Locations:** Geographic coordinates and addresses
- **User Data:** Customer and driver information
- **Business Metrics:** Revenue, ratings, performance indicators

Schema Design Challenge: How to efficiently store ride-booking data where each order has multiple events, locations, and associated metadata while supporting high-volume analytical queries?

3 Schema Design Approaches: Comprehensive Analysis

3.1 Approach 1: Normalized (Relational) Schema

Design Philosophy: Eliminate data redundancy through proper normalization

Implementation Strategy:

— *Normalized approach*

```
CREATE TABLE orders (
  order_id INT64 PRIMARY KEY,
  customer_id INT64,
  driver_id INT64,
  pickup_lat FLOAT64,
  pickup_lng FLOAT64,
  dropoff_lat FLOAT64,
  dropoff_lng FLOAT64,
  order_time TIMESTAMP,
  total_amount FLOAT64
);

CREATE TABLE order_events (
  event_id INT64 PRIMARY KEY,
  order_id INT64,
  event_type STRING,
  event_time TIMESTAMP,
  FOREIGN KEY (order_id) REFERENCES orders(order_id)
);
```

```
CREATE TABLE customers (  
    customer_id INT64 PRIMARY KEY,  
    name STRING,  
    phone STRING,  
    email STRING  
);
```

Performance Characteristics:

- **Storage Efficiency:** Minimal data redundancy
- **Data Integrity:** Strong referential constraints
- **Query Performance:** Expensive JOINS required
- **Scalability:** Poor performance with large datasets
- **Complexity:** Need to understand all table relationships

Analytical Query Example:

— *Complex query requiring multiple JOINS*

```
SELECT  
    o.order_id ,  
    c.name as customer_name ,  
    COUNT(e.event_id) as event_count ,  
    AVG(o.total_amount) as avg_amount  
FROM orders o  
JOIN customers c ON o.customer_id = c.customer_id  
JOIN order_events e ON o.order_id = e.order_id  
WHERE o.order_time >= '2024-01-01'  
GROUP BY o.order_id , c.name;
```

3.2 Approach 2: Fully Denormalized Schema

Design Philosophy: Pre-join all data for maximum query performance

Implementation Strategy:

— *Denormalized approach*

```
CREATE TABLE ride_data (  
    order_id INT64,  
    customer_id INT64,  
    customer_name STRING,  
    customer_phone STRING,  
    driver_id INT64,  
    driver_name STRING,  
    pickup_lat FLOAT64,  
    pickup_lng FLOAT64,  
    dropoff_lat FLOAT64,  
    dropoff_lng FLOAT64,  
    order_time TIMESTAMP,  
    event_type STRING,  
    event_time TIMESTAMP,
```

```
total_amount FLOAT64
);
```

Performance Characteristics:

- **Query Performance:** Fast single-table queries
- **Storage Overhead:** Significant data duplication
- **Aggregation Issues:** Risk of double-counting
- **Data Consistency:** Difficult to maintain
- **Update Complexity:** Changes require multiple row updates

Data Explosion Problem:

Order ID	Events	Rows Created
12345	4 events	4 rows (duplicate order data)
12346	3 events	3 rows (duplicate order data)
12347	5 events	5 rows (duplicate order data)

3.3 Approach 3: Nested and Repeated Fields (BigQuery Optimal)

Design Philosophy: Best of both worlds - logical organization with performance benefits

Implementation Strategy:

— *BigQuery nested and repeated approach*

```
CREATE TABLE orders (
  order_id INT64,
  customer STRUCT<
    id INT64,
    name STRING,
    phone STRING,
    email STRING
  >,
  driver STRUCT<
    id INT64,
    name STRING,
    rating FLOAT64
  >,
  pickup_location STRUCT<
    latitude FLOAT64,
    longitude FLOAT64,
    address STRING
  >,
  destination STRUCT<
    latitude FLOAT64,
    longitude FLOAT64,
    address STRING
  >,
  >
```

```

    events ARRAY<STRUCT<
        event_type STRING,
        timestamp TIMESTAMP,
        metadata STRUCT<
            location_lat FLOAT64,
            location_lng FLOAT64,
            status_code STRING
        >
    >>,
    order_time TIMESTAMP,
    total_amount FLOAT64,
    payment_method STRING
)
PARTITION BY DATE(order_time)
CLUSTER BY customer.id, driver.id;

```

Performance Characteristics:

- **Logical Organization:** One row per order
- **Query Performance:** No JOINS required
- **Storage Efficiency:** No data explosion
- **Data Integrity:** Maintains relationships
- **Analytical Power:** Supports complex aggregations

4 BigQuery Data Types Deep Dive

4.1 STRUCT Data Type: Nested Fields

Definition: Structured data type that allows grouping related fields within a single column

Key Characteristics:

- **Schema Type:** RECORD in BigQuery schema
- **Field Access:** Dot notation (e.g., customer.name)
- **Benefits:** Pre-joined data, logical organization
- **Use Cases:** Related data that should be co-located

Implementation Examples:

— *Simple STRUCT*

```

customer STRUCT<
    id INT64,
    name STRING,
    email STRING
>

```

— *Nested STRUCT*

```

address STRUCT<
  street STRUCT<
    number INT64,
    name STRING
  >,
  city STRING,
  state STRING,
  zip_code STRING
>

— STRUCT with arrays
contact_info STRUCT<
  phones ARRAY<STRING>,
  emails ARRAY<STRING>,
  primary_contact STRUCT<
    name STRING,
    phone STRING
  >
>

```

Query Examples:

```

— Accessing STRUCT fields
SELECT
  order_id ,
  customer.name ,
  customer.email ,
  pickup_location.latitude ,
  pickup_location.longitude
FROM orders
WHERE customer.name LIKE '%John%';

— Filtering on STRUCT fields
SELECT order_id , total_amount
FROM orders
WHERE pickup_location.latitude BETWEEN 40.0 AND 41.0
  AND pickup_location.longitude BETWEEN -74.0 AND -73.0;

```

4.2 ARRAY Data Type: Repeated Fields

Definition: Repeated values within a single field, allowing multiple values per row

Key Characteristics:

- **Mode:** REPEATED in BigQuery schema
- **Benefits:** Handles one-to-many relationships
- **Use Cases:** Events, timestamps, related items
- **Performance:** No JOIN operations needed

Implementation Examples:


```

— Simple ARRAY
tags ARRAY<STRING>

— ARRAY of STRUCTs
events ARRAY<STRUCT<
    event_type STRING,
    timestamp TIMESTAMP,
    user_id INT64
>>

— Nested ARRAYs
order_items ARRAY<STRUCT<
    product_id INT64,
    quantity INT64,
    price FLOAT64,
    options ARRAY<STRING>
>>

```

Query Examples:

```

— Working with ARRAYs
SELECT
    order_id ,
    ARRAY_LENGTH(events) as event_count ,
    events[OFFSET(0)].event_type as first_event
FROM orders;

```

```

— Unnesting ARRAYs
SELECT
    order_id ,
    event.event_type ,
    event.timestamp
FROM orders ,
    UNNEST(events) as event
WHERE event.timestamp >= '2024-01-01';

```

```

— Aggregating ARRAY data
SELECT
    order_id ,
    COUNT(*) as event_count ,
    ARRAY_AGG(event.event_type) as all_events
FROM orders ,
    UNNEST(events) as event
GROUP BY order_id;

```

5 Schema Design Decision Framework

5.1 When to Use Nested and Repeated Fields

Primary Indicators:

1. **One-to-Many Relationships:** Data naturally has parent-child structure

2. **Query Patterns:** Frequently query related data together
3. **Data Volume:** Large datasets where JOINS become expensive
4. **Analytical Workloads:** Complex aggregations and reporting
5. **Real-time Requirements:** Need for fast query response times

Specific Use Cases:

- **E-commerce:** Orders with multiple items and events
- **Analytics:** User sessions with multiple events
- **IoT Data:** Device readings with multiple sensors
- **Financial Data:** Transactions with multiple legs
- **Healthcare:** Patient records with multiple visits

5.2 When to Keep Normalized Schemas

Primary Indicators:

1. **Small Dimension Tables:** $\leq 10\text{GB}$ in size
2. **Frequent Updates:** High UPDATE/DELETE operations
3. **Complex Business Logic:** Multiple validation rules
4. **Legacy Integration:** Need to maintain compatibility
5. **Data Consistency:** Critical for business operations

Decision Matrix:

Factor	Use Nested/Repeated	Keep Normalized
Table Size	$\leq 10\text{GB}$	$\leq 10\text{GB}$
Update Frequency	Low	High
Query Complexity	Complex aggregations	Simple lookups
Data Relationships	One-to-many	Many-to-many
Performance Requirements	Critical	Moderate
Storage Cost Sensitivity	Low	High

6 Performance Optimization Strategies

6.1 Partitioning: Foundation of Performance

Concept: Dividing tables into smaller, manageable pieces based on column values

Benefits:

- **Cost Reduction:** Read only relevant partitions
- **Performance Improvement:** Faster query execution

- **Accurate Cost Estimation:** Better query planning
- **Automatic Management:** BigQuery handles partition creation

Partitioning Types:

1. **Date/Time Partitioning:** Most common and effective
2. **Integer Range Partitioning:** For ID-based partitioning
3. **Ingestion Time Partitioning:** Based on load time

Implementation Examples:

```

— Date partitioning
CREATE TABLE events (
  event_date DATE,
  user_id INT64,
  event_type STRING
)
PARTITION BY DATE(event_date);

— Integer range partitioning
CREATE TABLE user_data (
  user_id INT64,
  data STRING
)
PARTITION BY RANGE_BUCKET(user_id , GENERATE_ARRAY(0 , 1000000 , 10000));

— Partition with expiration
CREATE TABLE logs (
  log_date DATE,
  log_data STRING
)
PARTITION BY DATE(log_date)
OPTIONS(partition_expiration_days=90);

```

Best Practices:

- Always filter on partition columns in WHERE clauses
- Use date partitioning for time-series data
- Set appropriate expiration for old partitions
- Monitor partition count (metadata overhead)
- Isolate partition field on left side of comparisons

6.2 Clustering: Advanced Performance Tuning

Concept: Organizing data within partitions based on column values for optimal access patterns

Benefits:

- **Filter Performance:** Eliminates unnecessary data scans
- **Aggregation Performance:** Co-locates similar values
- **Cost Reduction:** Reduces bytes processed
- **Automatic Re-clustering:** BigQuery maintains optimal sorting

Implementation Examples:

— *Single column clustering*

```
CREATE TABLE user_events (
    event_date DATE,
    user_id INT64,
    event_type STRING
)
PARTITION BY DATE(event_date)
CLUSTER BY user_id;
```

— *Multi-column clustering*

```
CREATE TABLE orders (
    order_date DATE,
    customer_id INT64,
    product_category STRING,
    order_amount FLOAT64
)
PARTITION BY DATE(order_date)
CLUSTER BY customer_id, product_category;
```

— *Clustering with nested fields*

```
CREATE TABLE ride_data (
    ride_date DATE,
    customer STRUCT<id INT64, tier STRING>,
    driver STRUCT<id INT64, rating FLOAT64>,
    events ARRAY<STRUCT<type STRING, time TIMESTAMP>>
)
PARTITION BY DATE(ride_date)
CLUSTER BY customer.id, driver.id;
```

Clustering Best Practices:

- **Column Order Matters:** Most selective columns first
- **Combine with Partitioning:** Maximum performance benefit
- **Choose High-Cardinality Columns:** Better clustering effectiveness
- **Limit to 4 Columns:** Diminishing returns beyond this
- **Monitor Effectiveness:** Clustering weakens over time

6.3 Query Optimization Techniques

Schema-Level Optimization:

1. Use **nested/repeated fields** instead of JOINS when possible
2. **Keep dimension tables < 10GB** normalized
3. **Denormalize large tables** (> 10GB) for better performance
4. **Consider query patterns** when designing schema

Query-Level Optimization:

1. **Always filter on partition columns**
2. **Use clustering columns** in WHERE clauses
3. **Limit columns selected** to only what's needed
4. **Use appropriate data types** for better compression
5. **Avoid SELECT *** in large tables

Cost Optimization:

1. **Partition tables** for accurate cost estimation
2. **Use clustering** to reduce bytes processed
3. **Set partition expiration** for old data
4. **Monitor query costs** and optimize accordingly
5. **Use materialized views** for frequently accessed data

7 Real-World Implementation Examples

7.1 E-commerce Platform Schema

Business Requirements:

- **Scale:** Millions of orders daily
- **Data Types:** Orders, items, customers, payments
- **Queries:** Revenue analysis, customer behavior, inventory
- **Performance:** Sub-second response for dashboards

Optimized Schema Design:

```

CREATE TABLE orders (
  order_id INT64,
  customer STRUCT<
    id INT64,
    name STRING,
    email STRING,
    tier STRING,
    registration_date DATE
  >,
  order_info STRUCT<
    order_date DATE,
    status STRING,
    total_amount FLOAT64,
    tax_amount FLOAT64,
    shipping_amount FLOAT64
  >,
  items ARRAY<STRUCT<
    product_id INT64,
    product_name STRING,
    category STRING,
    quantity INT64,
    unit_price FLOAT64,
    total_price FLOAT64,
    options ARRAY<STRING>
  >>,
  payment STRUCT<
    method STRING,
    transaction_id STRING,
    status STRING,
    processed_at TIMESTAMP
  >,
  shipping STRUCT<
    address STRUCT<
      street STRING,
      city STRING,
      state STRING,
      zip_code STRING,
      country STRING
    >,
    method STRING,
    tracking_number STRING,
    estimated_delivery DATE
  >,
  events ARRAY<STRUCT<
    event_type STRING,
    timestamp TIMESTAMP,
    user_id INT64,
    metadata STRUCT<
      ip_address STRING,
      user_agent STRING,
      referrer STRING
    >
  >

```

```

    >
    >>
)
PARTITION BY DATE(order_info.order_date)
CLUSTER BY customer.id, customer.tier;

Query Examples:

— Revenue analysis by customer tier
SELECT
    customer.tier,
    COUNT(DISTINCT order_id) as order_count,
    SUM(order_info.total_amount) as total_revenue,
    AVG(order_info.total_amount) as avg_order_value
FROM orders
WHERE order_info.order_date >= '2024-01-01'
GROUP BY customer.tier;

— Product category analysis
SELECT
    item.category,
    COUNT(*) as item_count,
    SUM(item.total_price) as category_revenue
FROM orders,
    UNNEST(items) as item
WHERE order_info.order_date >= '2024-01-01'
GROUP BY item.category
ORDER BY category_revenue DESC;

— Customer journey analysis
SELECT
    customer.id,
    customer.name,
    ARRAY_LENGTH(events) as interaction_count,
    ARRAY_AGG(event.event_type) as event_sequence
FROM orders,
    UNNEST(events) as event
WHERE order_info.order_date >= '2024-01-01'
GROUP BY customer.id, customer.name;

```

7.2 Analytics Platform Schema

Business Requirements:

- **Scale:** Billions of events daily
- **Data Types:** User events, sessions, page views
- **Queries:** User behavior analysis, funnel analysis
- **Performance:** Real-time dashboards

Optimized Schema Design:

```

CREATE TABLE user_sessions (
  session_id STRING,
  user STRUCT<
    id INT64,
    email STRING,
    signup_date DATE,
    country STRING,
    device_type STRING
  >,
  session_info STRUCT<
    start_time TIMESTAMP,
    end_time TIMESTAMP,
    duration_seconds INT64,
    page_count INT64
  >,
  events ARRAY<STRUCT<
    event_type STRING,
    timestamp TIMESTAMP,
    page_url STRING,
    referrer STRING,
    properties STRUCT<
      button_clicked STRING,
      form_submitted BOOL,
      scroll_depth INT64,
      time_on_page INT64
    >
  >
  >>,
  conversion STRUCT<
    converted BOOL,
    conversion_type STRING,
    conversion_value FLOAT64,
    conversion_time TIMESTAMP
  >
)
PARTITION BY DATE(session_info.start_time)
CLUSTER BY user.id, user.country;

```

8 Interview Questions and Discussion Points

8.1 Core Concept Questions

1. **Q:** Explain the fundamental differences between normalized, denormalized, and nested schema designs in BigQuery. When would you choose each approach?

Expected Answer Points:

- Normalized: Eliminates redundancy, good for small tables, expensive JOINS
- Denormalized: Fast queries, data explosion, aggregation issues
- Nested: Best of both worlds, logical organization, no JOINS needed

- Decision factors: table size, update frequency, query patterns
2. **Q: How does BigQuery's columnar storage architecture impact schema design decisions?**

Expected Answer Points:

- Only reads required columns, ignores unused nested fields
 - 99 unused columns don't impact query performance
 - Pay only for data processed
 - Enables wide schemas without performance penalty
3. **Q: What are the performance implications of one-to-many relationships in different schema designs?**

Expected Answer Points:

- Normalized: Requires JOINS, becomes expensive at scale
- Denormalized: Data explosion, aggregation complexity
- Nested: No JOINS, maintains data integrity, efficient storage

8.2 Performance Optimization Questions

1. **Q: How would you optimize a BigQuery table for both filtering and aggregation queries?**

Expected Answer Points:

- Partition by date for time-based filtering
- Cluster by high-cardinality columns for equality filters
- Use nested fields to avoid JOINS
- Consider query patterns in schema design

2. **Q: When should you use partitioning vs clustering, and when should you use both?**

Expected Answer Points:

- Partitioning: For cost reduction and time-based queries
- Clustering: For filter performance and co-location
- Both: Maximum performance benefit, partition first then cluster
- Partitioning provides cost estimation, clustering provides performance

3. **Q: How do you handle schema evolution in BigQuery while maintaining performance?**

Expected Answer Points:

- Use backward-compatible schema changes
- Consider data migration strategies
- Monitor query performance after changes
- Use views for gradual transitions

8.3 Real-World Scenario Questions

1. **Q: Design a schema for a ride-sharing service processing millions of rides daily. Consider performance, cost, and analytical requirements.**

Expected Answer Points:

- Use nested fields for ride details, locations, events
- Partition by ride date for time-based queries
- Cluster by customer and driver IDs
- Include events array for ride lifecycle tracking
- Consider real-time vs batch processing requirements

2. **Q: How would you migrate a traditional relational schema to BigQuery? What are the key considerations?**

Expected Answer Points:

- Analyze query patterns and data relationships
- Identify opportunities for nested/repeated fields
- Plan partitioning and clustering strategy
- Consider data loading and transformation requirements
- Test performance with representative data volumes

3. **Q: Design a schema for a real-time analytics dashboard requiring sub-second response times.**

Expected Answer Points:

- Use materialized views for pre-computed aggregations
- Implement effective partitioning and clustering
- Consider data freshness vs performance trade-offs
- Use nested fields to minimize JOINS
- Optimize for specific dashboard queries

9 Advanced Topics and Best Practices

9.1 Schema Design Patterns

Star Schema Adaptation:

- **Fact Tables:** Use nested fields for dimension data
- **Dimension Tables:** Keep small dimensions normalized
- **Bridge Tables:** Use arrays for many-to-many relationships
- **Slowly Changing Dimensions:** Handle with nested structures

Data Vault Adaptation:

- **Hubs:** Use STRUCT for entity information
- **Links:** Use ARRAY for relationship data
- **Satellites:** Use nested fields for descriptive data
- **Point-in-Time Tables:** Use arrays for historical data

9.2 Performance Monitoring and Optimization

Key Metrics to Monitor:

- **Query Performance:** Execution time and slot utilization
- **Cost Efficiency:** Bytes processed and query costs
- **Partition Effectiveness:** Partition pruning statistics
- **Clustering Quality:** Data distribution and sorting

Optimization Strategies:

- **Regular Analysis:** Monitor query patterns and performance
- **Schema Refinement:** Adjust based on usage patterns
- **Partition Management:** Set appropriate expiration policies
- **Clustering Maintenance:** Monitor and adjust clustering columns

10 Key Takeaways and Best Practices

10.1 Schema Design Principles

- **Think in Arrays:** Use ARRAY data types for repeated values
- **Pre-join with STRUCTs:** Organize related data logically
- **Consider Query Patterns:** Design for your most common queries
- **Balance Flexibility and Performance:** Choose appropriate normalization level
- **Plan for Scale:** Design schemas that can handle growth

10.2 Performance Optimization

- **Partition by Date:** Most common and effective strategy
- **Cluster by High-Cardinality Columns:** Improve filter performance
- **Use Nested Fields:** Avoid expensive JOINS
- **Monitor and Optimize:** Continuously improve based on usage patterns
- **Consider Cost vs Performance:** Balance optimization with cost

10.3 Cost Management

- **Accurate Cost Estimation:** Use partitioning for better planning
- **Reduce Bytes Processed:** Use clustering and selective queries
- **Set Data Lifecycle:** Use partition expiration for cost control
- **Monitor Usage:** Track and optimize query patterns
- **Use Materialized Views:** Pre-compute expensive aggregations

10.4 Implementation Guidelines

- **Start with Partitioning:** Foundation for performance
- **Add Clustering Strategically:** Based on query patterns
- **Use Nested Fields Judiciously:** When relationships are clear
- **Test with Real Data:** Validate performance assumptions
- **Document Design Decisions:** For future maintenance